

Residual-Based Adaptive Sampling for PINNs

Paper Reference

Title: A Comprehensive Study of Non-adaptive and Residual-based Adaptive Sampling for Physics-informed Neural Networks

Authors: Wu et al. (2022)

Journal: Computer Methods in Applied Mechanics and Engineering

DOI: [10.1016/j.cma.2022.115671](https://doi.org/10.1016/j.cma.2022.115671)

Code Repository: <https://github.com/lu-group/pinn-sampling>

Overview

This paper systematically compares **10 different sampling methods** for selecting collocation (residual) points in Physics-Informed Neural Networks (PINNs). The study includes over **6,000 PINN simulations** across 6 different PDE problems.

Sampling Methods

A. Non-Adaptive Uniform Sampling (6 methods)

These methods select points without considering the PDE residual:

Method	Description	Implementation Complexity
Grid	Regular equispaced grid points	Very Easy
Random	Uniformly random sampling	Very Easy
LHS	Latin Hypercube Sampling - stratified random	Easy (scipy)
Halton	Low-discrepancy quasi-random sequence	Easy (scipy)
Hammersley	Quasi-random sequence	Easy (scipy)
Sobol	Low-discrepancy sequence	Easy (scipy)

B. Uniform with Resampling (Random-R)

- Same as Random sampling, but **resample points every N iterations**
- The resampling period **N** is an important hyperparameter
- Special case of RAD with **k=0** or **C→∞**

C. Adaptive Nonuniform Sampling (3 methods)

These methods use the **PDE residual** to guide point selection:

1. RAR-G (Residual-based Adaptive Refinement with Greed) - Algorithm 1

- Original method from Lu et al. (2019)
- After every certain number of iterations, **adds new points** only in locations with the **largest PDE residuals**
- Greedy approach - focuses only on maximum residual regions
- Points **accumulate** over training (dataset grows)
- Special case of RAR-D with $k \rightarrow \infty$

Algorithm 1: RAR-G

1. Initialize residual points S randomly
2. Train PINN for some iterations
3. Evaluate residual $\varepsilon(x)$ on dense candidate set
4. Add points with largest residuals to S
5. Repeat from step 2

2. RAD (Residual-based Adaptive Distribution) - Algorithm 2

- **Resamples ALL points** according to a probability density function (PDF)
- Does NOT accumulate points - replaces entire training set
- PDF formula:

$$\text{p}(\mathbf{x}) \propto \frac{\varepsilon^k(\mathbf{x})}{E[\varepsilon^k(\mathbf{x})]} + c$$

Where:

- $\varepsilon(\mathbf{x}) = f(\mathbf{x}; \hat{u}(\mathbf{x}))$ is the PDE residual
- $k \geq 0$ controls how much to focus on high-residual regions
- $c \geq 0$ ensures some uniform coverage (regularization)

Special cases:

Parameters	Equivalent Method
$k=0$ or $c \rightarrow \infty$	Random-R (uniform resampling)
$k=1, c=0$	Nabian et al.'s method
$k=2, c=0$	Recommended default

Algorithm 2: RAD

1. Initialize residual points S randomly
2. Train PINN for some iterations
3. Evaluate residual $\varepsilon(x)$ on dense candidate set S_0
4. Compute probability $p(x) \propto \varepsilon^k(x)/E[\varepsilon^k(x)] + c$
5. Resample ALL points from S_0 according to $p(x)$
6. Repeat from step 2

3. RAR-D (Residual-based Adaptive Refinement with Distribution) - Algorithm 3

- **Hybrid** of RAR-G and RAD
- **Repeatedly adds new points** (like RAR-G)
- But new points are **sampled according to the RAD PDF** (not just max residual)
- Points accumulate over training

Special cases:

Parameters	Equivalent Method
$k \rightarrow \infty$	RAR-G (greedy)
$k=2, c=0$	Recommended default

Algorithm 3: RAR-D

1. Initialize residual points S randomly
2. Train PINN for some iterations
3. Evaluate residual $\varepsilon(x)$ on dense candidate set S_0
4. Compute probability $p(x) \propto \varepsilon^k(x)/E[\varepsilon^k(x)] + c$
5. Sample M new points from S_0 according to $p(x)$
6. Add new points to S (S grows)
7. Repeat from step 2

Benchmark PDEs

The paper tests all methods on **6 PDEs** (4 forward + 2 inverse problems):

Forward Problems

Problem	PDE	Domain	Network	Optimizer
Diffusion	$\frac{\partial u}{\partial t} = \frac{\partial^2 u}{\partial x^2} + f(x, t)$	$x \in [-1, 1], t \in [0, 1]$	4×32	Adam
Burgers'	$\frac{\partial u}{\partial t} + u \frac{\partial u}{\partial x} = \nu \frac{\partial^2 u}{\partial x^2}$	$x \in [-1, 1], t \in [0, 1]$	4×64	Adam + L-BFGS
Allen-Cahn	$\frac{\partial u}{\partial t} = D \frac{\partial^2 u}{\partial x^2} + 5(u - u^3)$	$x \in [-1, 1], t \in [0, 1]$	4×64	Adam + L-BFGS
Wave (multi-scale)	$\frac{\partial^2 u}{\partial t^2} = 4 \frac{\partial^2 u}{\partial x^2}$	$x \in [0, 1], t \in [0, 1]$	6×100	Adam + L-BFGS

Inverse Problems

Problem	PDE	Parameters to Infer	Network	Optimizer
---------	-----	---------------------	---------	-----------

Problem	PDE	Parameters to Infer	Network	Optimizer
Diffusion-Reaction	$\frac{\partial u}{\partial t} = D \frac{\partial^2 u}{\partial x^2} + ku^2$	\$D\$, \$k\$	4x20	Adam
Korteweg-de Vries (KdV)	$\frac{\partial u}{\partial t} + \lambda_1 u \frac{\partial u}{\partial x} + \lambda_2 \frac{\partial^3 u}{\partial x^3} = 0$	λ_1 , λ_2	4x100	Adam

Key Experimental Results

Forward Problems (L^2 Relative Error)

Method	Diffusion	Burgers'	Allen-Cahn	Wave
Grid	0.66%	13.7%	93.4%	81.3%
Random	0.74%	13.3%	22.2%	68.4%
LHS	0.48%	13.5%	26.6%	75.9%
Halton	0.24%	4.51%	0.29%	60.2%
Hammersley	0.17%	3.02%	0.14%	58.9%
Sobol	0.19%	3.38%	0.35%	57.5%
Random-R	0.12%	1.69%	0.55%	0.72%
RAR-G	0.20%	0.12%	0.53%	0.81%
RAD	0.11%	0.02%	0.08%	0.09%
RAR-D	0.14%	0.03%	0.09%	0.29%

Bold = top 3 performers

Key Findings

- RAD performs best overall** among all 10 sampling methods for both forward and inverse problems
- For PDEs with complicated solutions** (Burgers', multi-scale Wave):
 - RAD and RAR-D are **dramatically better** (errors magnitudes lower)
 - 0.02% vs 13% for Burgers' (650x improvement!)
- For PDEs with smooth solutions** (Diffusion, Diffusion-Reaction):
 - Some uniform methods (Hammersley, Random-R) also perform well
 - Adaptive methods still best, but margin is smaller
- Low-discrepancy sequences** (Halton, Hammersley, Sobol) generally outperform Random and LHS

5. **Resampling helps:** Random-R consistently beats fixed Random sampling

6. Recommended hyperparameters:

- RAD: `k=2, c=0` (default)
 - RAR-D: `k=2, c=0` (default)
 - Resampling period: problem-dependent, typically every 100-1000 iterations
-

Implementation Recommendations for pinn-point

Easy to Implement (High Priority)

1. **Random-R:** Simply resample uniform random points every N epochs

- Minimal code change from current random sampling
- Good baseline improvement

2. **RAD:** At each resampling step:

- Evaluate residuals on dense candidate set
- Compute weights: `w = residual^k / mean(residual^k) + c`
- Sample points proportionally to weights
- Replace training set

3. **Low-discrepancy sequences:** Use scipy's `qmc` module

```
from scipy.stats import qmc
sampler = qmc.Halton(d=2) # or Sobol, LatinHypercube
points = sampler.random(n=1000)
```

Moderate Complexity

4. **RAR-D:** Similar to RAD, but:

- Keep existing points
- Add M new points sampled by PDF
- Dataset grows over time

5. **RAR-G:** Like current adaptive refinement, but:

- Use top-k residual points instead of mesh refinement
 - Add points greedily
-

Comparison with Current pinn-point Approach

The current `pinn-point` adaptive method uses **mesh refinement** based on FEM error estimation, which is conceptually similar to **RAR-G** but:

Aspect	plnn-point (current)	RAR-G (paper)
Refinement basis	FEM mesh + error estimator	PDE residual magnitude
Point placement	Mesh vertices	Any location (meshless)
Geometry	Complex geometries via NGSolve	Simple domains
Computational cost	Higher (mesh operations)	Lower

The paper's methods are **meshless** and potentially simpler to implement, but our mesh-based approach handles complex geometries naturally.

References

- [1] Wu et al. (2022) - This paper
- [2] Lu et al. (2019) - Original RAR method (DeepXDE)
- [3] Nabian et al. (2021) - PDF-based resampling